

```
In [1]: import numpy as np
import pandas as pd
from keras.models import Sequential
from keras.layers import Conv2D, MaxPooling2D, Activation, Dropout, Flatten, Dense
from keras.regularizers import l2
from tensorflow.keras.optimizers import Adam
from keras.callbacks import ReduceLROnPlateau
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
import matplotlib.pyplot as plt
```

```
In [8]: # Load and preprocess the data
import numpy as np
import json

def load_data(train_file_path, test_file_path):
    """
        Load training and test data from JSON files and return as arrays.

    Parameters:
    train_file_path (str): Path to the training JSON file.
    test_file_path (str): Path to the test JSON file.

    Returns:
    tuple: X_train, y_train, X_test arrays
    """
    # Load the JSON data
    with open(train_file_path) as train_file:
        train_data = json.load(train_file)
    with open(test_file_path) as test_file:
        test_data = json.load(test_file)

    # Extract features and labels, convert to numpy arrays
    X_train = np.array([np.array(instance['band_1'], dtype=np.float32).reshape(75, 75) for instance in train_data])
    y_train = np.array([instance['is_iceberg'] for instance in train_data], dtype=np.int32)

    return X_train, y_train, X_test

# Usage example:
# X_train, y_train, X_test = load_data('path/to/train.json', 'path/to/test.json')

def normalize(X):
    return (X - np.min(X)) / (np.max(X) - np.min(X))

def preprocess(X):
    X_normalized = np.array([normalize(image) for image in X])
    return X_normalized.reshape(-1, 75, 75, 1)

X_train, y_train, X_test = load_data(r'train.json\data\processed\train.json', r'test.json')
X_train = preprocess(X_train)
X_test = preprocess(X_test)

# Train-test split for validation
X_train, X_val, y_train, y_val = train_test_split(X_train, y_train, test_size=0.2)
```

```
In [9]: # Define model hyperparameters
batch_size = 32
num_epochs = 50
kernel_size = 3
pool_size = 2
conv_depth_1 = 32
conv_depth_2 = 64
conv_depth_3 = 128
dense_1 = 128
dense_2 = 1
drop_out = 0.3
weight_decay = 1e-4

# Model architecture with modifications
model = Sequential()
model.add(Conv2D(conv_depth_1, (kernel_size, kernel_size), input_shape=(75, 75,
model.add(BatchNormalization())
model.add(Activation('relu')))
model.add(MaxPooling2D(pool_size=(pool_size, pool_size)))

model.add(Conv2D(conv_depth_2, (kernel_size, kernel_size), padding='same', kerne
model.add(BatchNormalization())
model.add(Activation('relu')))
model.add(MaxPooling2D(pool_size=(pool_size, pool_size)))

model.add(Conv2D(conv_depth_3, (kernel_size, kernel_size), padding='same', kerne
model.add(BatchNormalization())
model.add(Activation('relu')))
model.add(MaxPooling2D(pool_size=(pool_size, pool_size)))

model.add(Flatten())
model.add(Dense(dense_1, kernel_regularizer=l2(weight_decay)))
model.add(Activation('relu'))
model.add(Dropout(drop_out))

model.add(Dense(dense_2, activation='sigmoid'))

# Compile the model
optimizer = Adam(learning_rate=0.001)
reduce_lr = ReduceLROnPlateau(monitor='val_loss', factor=0.5, patience=3, min_lr
model.compile(optimizer=optimizer, loss='binary_crossentropy', metrics=['accurac

c:\Users\sanba\anaconda3\Lib\site-packages\keras\src\layers\convolutional\base_co
nv.py:107: UserWarning: Do not pass an `input_shape`/`input_dim` argument to a la
yer. When using Sequential models, prefer using an `Input(shape)` object as the f
irst layer in the model instead.
super().__init__(activity_regularizer=activity_regularizer, **kwargs)
```

```
In [10]: # Model training
history = model.fit(
    X_train, y_train,
    batch_size=batch_size,
    epochs=num_epochs,
    validation_data=(X_val, y_val),
    callbacks=[reduce_lr]
)

# Save model
model.save_weights('modified_model.weights.h5')
```

Epoch 1/50
41/41 19s 307ms/step - accuracy: 0.5809 - loss: 2.5957 - val_accuracy: 0.4953 - val_loss: 0.7304 - learning_rate: 0.0010
Epoch 2/50
41/41 12s 293ms/step - accuracy: 0.6657 - loss: 0.6402 - val_accuracy: 0.4891 - val_loss: 0.7416 - learning_rate: 0.0010
Epoch 3/50
41/41 11s 261ms/step - accuracy: 0.7124 - loss: 0.5966 - val_accuracy: 0.4891 - val_loss: 0.7706 - learning_rate: 0.0010
Epoch 4/50
41/41 11s 255ms/step - accuracy: 0.7307 - loss: 0.5493 - val_accuracy: 0.4891 - val_loss: 0.9636 - learning_rate: 0.0010
Epoch 5/50
41/41 11s 260ms/step - accuracy: 0.7634 - loss: 0.5021 - val_accuracy: 0.4891 - val_loss: 0.8819 - learning_rate: 5.0000e-04
Epoch 6/50
41/41 11s 257ms/step - accuracy: 0.8102 - loss: 0.4660 - val_accuracy: 0.4891 - val_loss: 0.9208 - learning_rate: 5.0000e-04
Epoch 7/50
41/41 11s 256ms/step - accuracy: 0.7970 - loss: 0.4819 - val_accuracy: 0.4891 - val_loss: 0.9537 - learning_rate: 5.0000e-04
Epoch 8/50
41/41 10s 249ms/step - accuracy: 0.8278 - loss: 0.4253 - val_accuracy: 0.4891 - val_loss: 0.8965 - learning_rate: 2.5000e-04
Epoch 9/50
41/41 10s 251ms/step - accuracy: 0.8215 - loss: 0.4527 - val_accuracy: 0.4891 - val_loss: 0.8936 - learning_rate: 2.5000e-04
Epoch 10/50
41/41 11s 255ms/step - accuracy: 0.8472 - loss: 0.4222 - val_accuracy: 0.4891 - val_loss: 0.9120 - learning_rate: 2.5000e-04
Epoch 11/50
41/41 10s 250ms/step - accuracy: 0.8592 - loss: 0.4046 - val_accuracy: 0.4953 - val_loss: 0.8982 - learning_rate: 1.2500e-04
Epoch 12/50
41/41 11s 258ms/step - accuracy: 0.8625 - loss: 0.4136 - val_accuracy: 0.5421 - val_loss: 0.8499 - learning_rate: 1.2500e-04
Epoch 13/50
41/41 12s 303ms/step - accuracy: 0.8614 - loss: 0.3870 - val_accuracy: 0.5452 - val_loss: 0.8925 - learning_rate: 1.2500e-04
Epoch 14/50
41/41 11s 278ms/step - accuracy: 0.8830 - loss: 0.3640 - val_accuracy: 0.5514 - val_loss: 0.9141 - learning_rate: 6.2500e-05
Epoch 15/50
41/41 10s 246ms/step - accuracy: 0.8799 - loss: 0.3885 - val_accuracy: 0.6604 - val_loss: 0.6872 - learning_rate: 6.2500e-05
Epoch 16/50
41/41 9s 221ms/step - accuracy: 0.8674 - loss: 0.3714 - val_accuracy: 0.6698 - val_loss: 0.6683 - learning_rate: 6.2500e-05
Epoch 17/50
41/41 9s 226ms/step - accuracy: 0.8797 - loss: 0.3592 - val_accuracy: 0.6854 - val_loss: 0.6531 - learning_rate: 6.2500e-05
Epoch 18/50
41/41 9s 221ms/step - accuracy: 0.8670 - loss: 0.3895 - val_accuracy: 0.7445 - val_loss: 0.5261 - learning_rate: 6.2500e-05
Epoch 19/50
41/41 9s 216ms/step - accuracy: 0.8800 - loss: 0.3628 - val_accuracy: 0.7695 - val_loss: 0.4718 - learning_rate: 6.2500e-05
Epoch 20/50
41/41 9s 214ms/step - accuracy: 0.8748 - loss: 0.3618 - val_accuracy: 0.8349 - val_loss: 0.3978 - learning_rate: 6.2500e-05

Epoch 21/50
41/41 9s 214ms/step - accuracy: 0.8808 - loss: 0.3688 - val_accuracy: 0.8505 - val_loss: 0.3838 - learning_rate: 6.2500e-05
Epoch 22/50
41/41 9s 213ms/step - accuracy: 0.8879 - loss: 0.3551 - val_accuracy: 0.8193 - val_loss: 0.4141 - learning_rate: 6.2500e-05
Epoch 23/50
41/41 9s 214ms/step - accuracy: 0.8819 - loss: 0.3481 - val_accuracy: 0.8224 - val_loss: 0.3944 - learning_rate: 6.2500e-05
Epoch 24/50
41/41 9s 216ms/step - accuracy: 0.8686 - loss: 0.3691 - val_accuracy: 0.8287 - val_loss: 0.3848 - learning_rate: 6.2500e-05
Epoch 25/50
41/41 9s 217ms/step - accuracy: 0.8695 - loss: 0.3678 - val_accuracy: 0.8442 - val_loss: 0.3786 - learning_rate: 3.1250e-05
Epoch 26/50
41/41 9s 218ms/step - accuracy: 0.8888 - loss: 0.3675 - val_accuracy: 0.8505 - val_loss: 0.3774 - learning_rate: 3.1250e-05
Epoch 27/50
41/41 9s 212ms/step - accuracy: 0.8780 - loss: 0.3507 - val_accuracy: 0.8442 - val_loss: 0.3751 - learning_rate: 3.1250e-05
Epoch 28/50
41/41 9s 212ms/step - accuracy: 0.8802 - loss: 0.3694 - val_accuracy: 0.8505 - val_loss: 0.3740 - learning_rate: 3.1250e-05
Epoch 29/50
41/41 9s 213ms/step - accuracy: 0.8965 - loss: 0.3256 - val_accuracy: 0.8505 - val_loss: 0.3745 - learning_rate: 3.1250e-05
Epoch 30/50
41/41 9s 213ms/step - accuracy: 0.8993 - loss: 0.3300 - val_accuracy: 0.8318 - val_loss: 0.3975 - learning_rate: 3.1250e-05
Epoch 31/50
41/41 9s 212ms/step - accuracy: 0.8880 - loss: 0.3651 - val_accuracy: 0.8567 - val_loss: 0.3713 - learning_rate: 3.1250e-05
Epoch 32/50
41/41 9s 210ms/step - accuracy: 0.9058 - loss: 0.3451 - val_accuracy: 0.8224 - val_loss: 0.3856 - learning_rate: 3.1250e-05
Epoch 33/50
41/41 9s 215ms/step - accuracy: 0.8714 - loss: 0.3706 - val_accuracy: 0.8567 - val_loss: 0.3698 - learning_rate: 3.1250e-05
Epoch 34/50
41/41 10s 203ms/step - accuracy: 0.8937 - loss: 0.3483 - val_accuracy: 0.8349 - val_loss: 0.3874 - learning_rate: 3.1250e-05
Epoch 35/50
41/41 10s 250ms/step - accuracy: 0.8826 - loss: 0.3382 - val_accuracy: 0.8287 - val_loss: 0.3758 - learning_rate: 3.1250e-05
Epoch 36/50
41/41 9s 229ms/step - accuracy: 0.8905 - loss: 0.3505 - val_accuracy: 0.8380 - val_loss: 0.3755 - learning_rate: 3.1250e-05
Epoch 37/50
41/41 9s 214ms/step - accuracy: 0.9012 - loss: 0.3467 - val_accuracy: 0.8567 - val_loss: 0.3691 - learning_rate: 1.5625e-05
Epoch 38/50
41/41 9s 211ms/step - accuracy: 0.8974 - loss: 0.3294 - val_accuracy: 0.8567 - val_loss: 0.3697 - learning_rate: 1.5625e-05
Epoch 39/50
41/41 9s 210ms/step - accuracy: 0.9040 - loss: 0.3137 - val_accuracy: 0.8536 - val_loss: 0.3695 - learning_rate: 1.5625e-05
Epoch 40/50
41/41 9s 214ms/step - accuracy: 0.9012 - loss: 0.3278 - val_accuracy: 0.8536 - val_loss: 0.3683 - learning_rate: 1.5625e-05

```

Epoch 41/50
41/41 9s 211ms/step - accuracy: 0.9082 - loss: 0.3288 - val_
accuracy: 0.8536 - val_loss: 0.3689 - learning_rate: 1.5625e-05
Epoch 42/50
41/41 9s 209ms/step - accuracy: 0.9090 - loss: 0.3318 - val_
accuracy: 0.8598 - val_loss: 0.3652 - learning_rate: 1.5625e-05
Epoch 43/50
41/41 9s 210ms/step - accuracy: 0.9058 - loss: 0.3325 - val_
accuracy: 0.8536 - val_loss: 0.3668 - learning_rate: 1.5625e-05
Epoch 44/50
41/41 9s 213ms/step - accuracy: 0.9024 - loss: 0.3213 - val_
accuracy: 0.8442 - val_loss: 0.3696 - learning_rate: 1.5625e-05
Epoch 45/50
41/41 9s 217ms/step - accuracy: 0.8936 - loss: 0.3344 - val_
accuracy: 0.8474 - val_loss: 0.3685 - learning_rate: 1.5625e-05
Epoch 46/50
41/41 9s 211ms/step - accuracy: 0.9141 - loss: 0.3321 - val_
accuracy: 0.8629 - val_loss: 0.3672 - learning_rate: 1.0000e-05
Epoch 47/50
41/41 10s 244ms/step - accuracy: 0.9119 - loss: 0.3144 - val_
accuracy: 0.8505 - val_loss: 0.3674 - learning_rate: 1.0000e-05
Epoch 48/50
41/41 9s 221ms/step - accuracy: 0.9129 - loss: 0.3305 - val_
accuracy: 0.8629 - val_loss: 0.3662 - learning_rate: 1.0000e-05
Epoch 49/50
41/41 9s 215ms/step - accuracy: 0.9176 - loss: 0.3091 - val_
accuracy: 0.8598 - val_loss: 0.3653 - learning_rate: 1.0000e-05
Epoch 50/50
41/41 9s 215ms/step - accuracy: 0.8897 - loss: 0.3483 - val_
accuracy: 0.8536 - val_loss: 0.3673 - learning_rate: 1.0000e-05

```

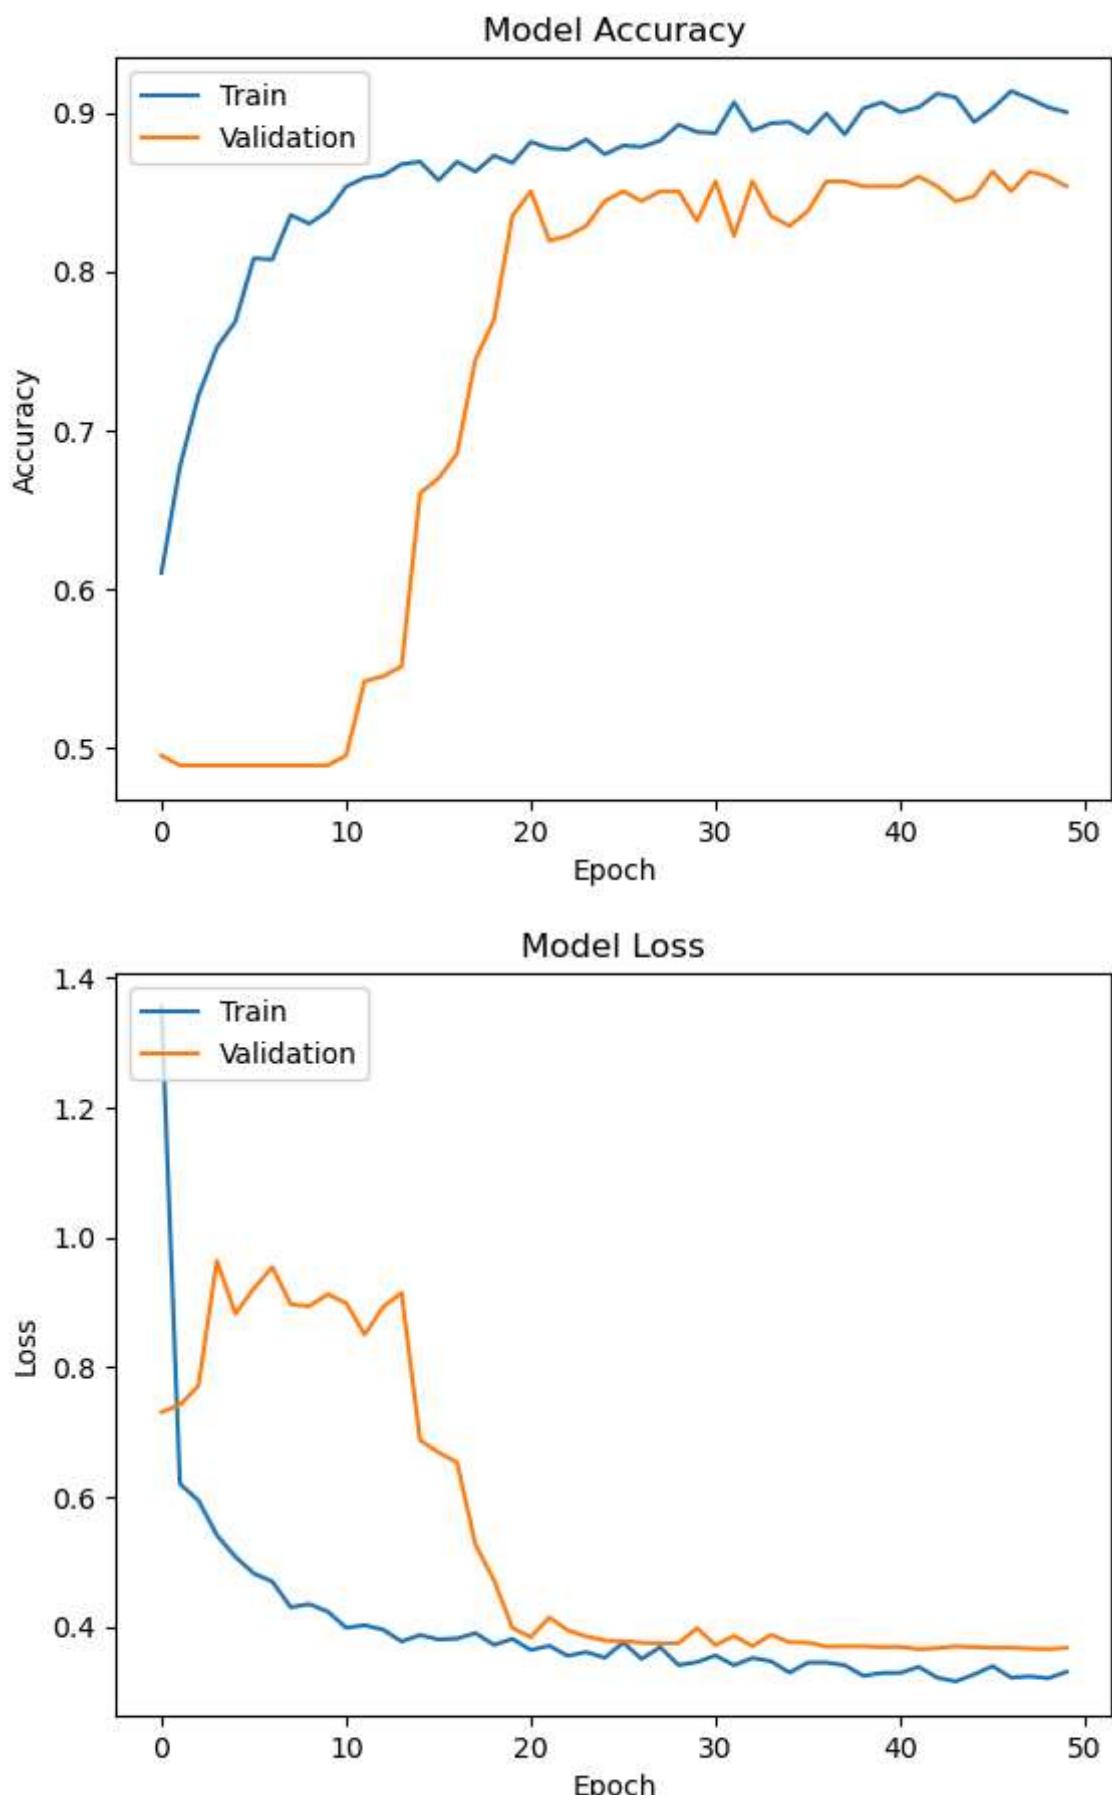
```

In [11]: # Plotting accuracy
plt.plot(history.history['accuracy'])
plt.plot(history.history['val_accuracy'])
plt.title('Model Accuracy')
plt.ylabel('Accuracy')
plt.xlabel('Epoch')
plt.legend(['Train', 'Validation'], loc='upper left')
plt.show()

# Plotting loss
plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])
plt.title('Model Loss')
plt.ylabel('Loss')
plt.xlabel('Epoch')
plt.legend(['Train', 'Validation'], loc='upper left')
plt.show()

# Prediction example (optional)
predictions = model.predict(X_test)
print(predictions[:5])

```



264/264 ━━━━━━ 13s 49ms/step

```
[[0.10773922]
 [0.80737066]
 [0.18277249]
 [0.9919188 ]
 [0.94463503]]
```

```
In [12]: # Train-test split for validation
X_train, X_val, y_train, y_val = train_test_split(X_train, y_train, test_size=0.

In [13]: # Define model hyperparameters
batch_size = 32
num_epochs = 50
kernel_size = 3
pool_size = 2
conv_depth_1 = 32
conv_depth_2 = 64
conv_depth_3 = 128
dense_1 = 128
dense_2 = 1
drop_out = 0.3
weight_decay = 1e-4

# Model architecture with modifications
model = Sequential()
model.add(Conv2D(conv_depth_1, (kernel_size, kernel_size), input_shape=(75, 75,
model.add(BatchNormalization())
model.add(Activation('relu')))
model.add(MaxPooling2D(pool_size=(pool_size, pool_size)))

model.add(Conv2D(conv_depth_2, (kernel_size, kernel_size), padding='same', kerne
model.add(BatchNormalization())
model.add(Activation('relu')))
model.add(MaxPooling2D(pool_size=(pool_size, pool_size)))

model.add(Conv2D(conv_depth_3, (kernel_size, kernel_size), padding='same', kerne
model.add(BatchNormalization())
model.add(Activation('relu')))
model.add(MaxPooling2D(pool_size=(pool_size, pool_size)))

model.add(Flatten())
model.add(Dense(dense_1, kernel_regularizer=l2(weight_decay)))
model.add(Activation('relu'))
model.add(Dropout(drop_out))

model.add(Dense(dense_2, activation='sigmoid'))

# Compile the model
optimizer = Adam(learning_rate=0.001)
reduce_lr = ReduceLROnPlateau(monitor='val_loss', factor=0.5, patience=3, min_lr
model.compile(optimizer=optimizer, loss='binary_crossentropy', metrics=['accuracy'])

# Model training
history = model.fit(
    X_train, y_train,
    batch_size=batch_size,
    epochs=num_epochs,
    validation_data=(X_val, y_val),
    callbacks=[reduce_lr]
)

# Save model
model.save_weights('modified_model.weights.h5')

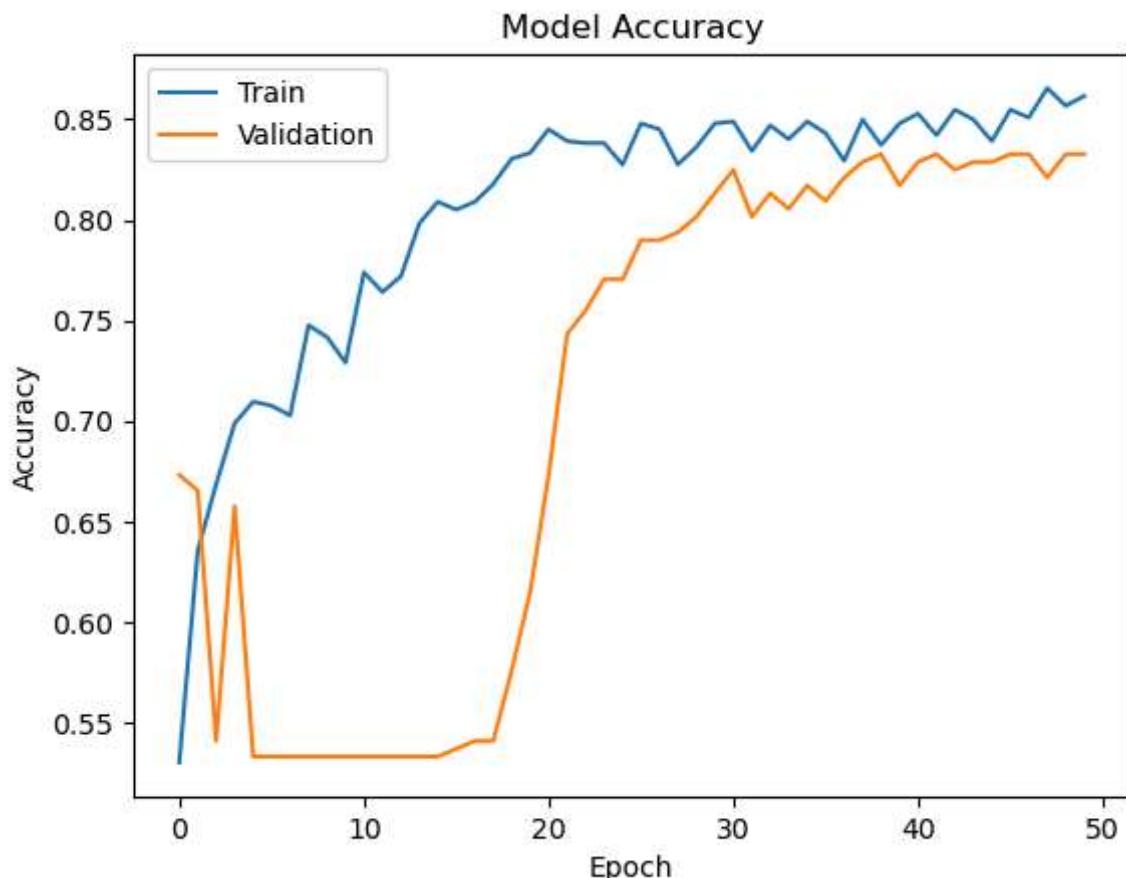
# Plotting accuracy
plt.plot(history.history['accuracy'])
```

```
plt.plot(history.history['val_accuracy'])
plt.title('Model Accuracy')
plt.ylabel('Accuracy')
plt.xlabel('Epoch')
plt.legend(['Train', 'Validation'], loc='upper left')
plt.show()
```

Epoch 1/50
33/33 12s 228ms/step - accuracy: 0.4999 - loss: 5.4829 - val_accuracy: 0.6732 - val_loss: 0.7146 - learning_rate: 0.0010
Epoch 2/50
33/33 7s 221ms/step - accuracy: 0.6245 - loss: 0.9461 - val_accuracy: 0.6654 - val_loss: 0.7249 - learning_rate: 0.0010
Epoch 3/50
33/33 8s 239ms/step - accuracy: 0.6794 - loss: 0.6412 - val_accuracy: 0.5409 - val_loss: 0.7152 - learning_rate: 0.0010
Epoch 4/50
33/33 7s 214ms/step - accuracy: 0.6942 - loss: 0.5893 - val_accuracy: 0.6576 - val_loss: 0.7132 - learning_rate: 0.0010
Epoch 5/50
33/33 7s 216ms/step - accuracy: 0.6945 - loss: 0.5942 - val_accuracy: 0.5331 - val_loss: 0.7401 - learning_rate: 0.0010
Epoch 6/50
33/33 7s 210ms/step - accuracy: 0.7241 - loss: 0.5488 - val_accuracy: 0.5331 - val_loss: 0.8211 - learning_rate: 0.0010
Epoch 7/50
33/33 7s 214ms/step - accuracy: 0.7236 - loss: 0.5523 - val_accuracy: 0.5331 - val_loss: 0.8690 - learning_rate: 0.0010
Epoch 8/50
33/33 7s 211ms/step - accuracy: 0.7406 - loss: 0.5589 - val_accuracy: 0.5331 - val_loss: 0.9033 - learning_rate: 5.0000e-04
Epoch 9/50
33/33 7s 219ms/step - accuracy: 0.7429 - loss: 0.5463 - val_accuracy: 0.5331 - val_loss: 0.8189 - learning_rate: 5.0000e-04
Epoch 10/50
33/33 7s 210ms/step - accuracy: 0.7288 - loss: 0.5543 - val_accuracy: 0.5331 - val_loss: 0.8465 - learning_rate: 5.0000e-04
Epoch 11/50
33/33 7s 209ms/step - accuracy: 0.7785 - loss: 0.5199 - val_accuracy: 0.5331 - val_loss: 0.9286 - learning_rate: 2.5000e-04
Epoch 12/50
33/33 7s 209ms/step - accuracy: 0.7597 - loss: 0.4948 - val_accuracy: 0.5331 - val_loss: 1.1116 - learning_rate: 2.5000e-04
Epoch 13/50
33/33 7s 211ms/step - accuracy: 0.7561 - loss: 0.4937 - val_accuracy: 0.5331 - val_loss: 1.2812 - learning_rate: 2.5000e-04
Epoch 14/50
33/33 7s 206ms/step - accuracy: 0.7825 - loss: 0.4434 - val_accuracy: 0.5331 - val_loss: 1.3986 - learning_rate: 1.2500e-04
Epoch 15/50
33/33 7s 215ms/step - accuracy: 0.8095 - loss: 0.4420 - val_accuracy: 0.5331 - val_loss: 1.4874 - learning_rate: 1.2500e-04
Epoch 16/50
33/33 7s 215ms/step - accuracy: 0.7988 - loss: 0.4576 - val_accuracy: 0.5370 - val_loss: 1.3904 - learning_rate: 1.2500e-04
Epoch 17/50
33/33 7s 208ms/step - accuracy: 0.7807 - loss: 0.4528 - val_accuracy: 0.5409 - val_loss: 1.2923 - learning_rate: 6.2500e-05
Epoch 18/50
33/33 7s 208ms/step - accuracy: 0.8186 - loss: 0.4230 - val_accuracy: 0.5409 - val_loss: 1.2340 - learning_rate: 6.2500e-05
Epoch 19/50
33/33 7s 209ms/step - accuracy: 0.8132 - loss: 0.4369 - val_accuracy: 0.5759 - val_loss: 1.0517 - learning_rate: 6.2500e-05
Epoch 20/50
33/33 7s 208ms/step - accuracy: 0.8193 - loss: 0.4265 - val_accuracy: 0.6148 - val_loss: 0.9028 - learning_rate: 3.1250e-05

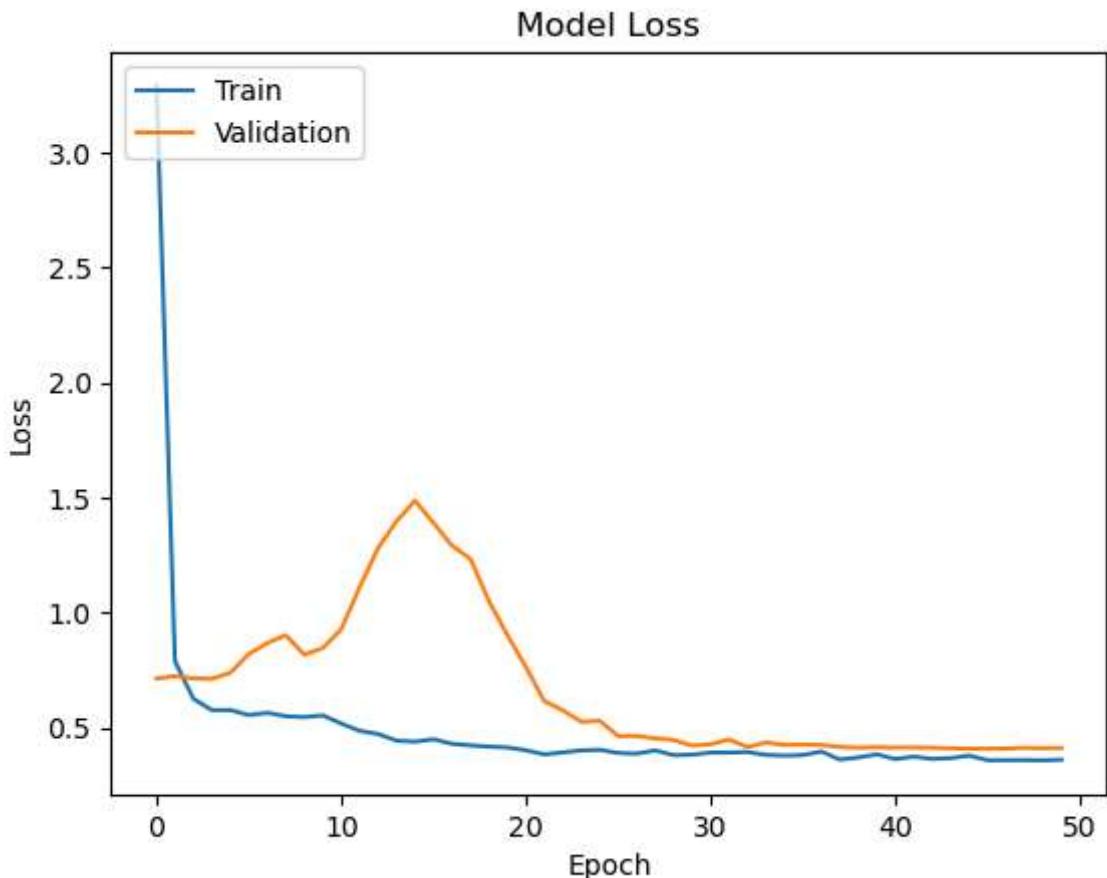
Epoch 21/50
33/33 7s 208ms/step - accuracy: 0.8369 - loss: 0.3972 - val_accuracy: 0.6732 - val_loss: 0.7651 - learning_rate: 3.1250e-05
Epoch 22/50
33/33 7s 210ms/step - accuracy: 0.8383 - loss: 0.3802 - val_accuracy: 0.7432 - val_loss: 0.6153 - learning_rate: 3.1250e-05
Epoch 23/50
33/33 7s 211ms/step - accuracy: 0.8387 - loss: 0.3938 - val_accuracy: 0.7549 - val_loss: 0.5771 - learning_rate: 3.1250e-05
Epoch 24/50
33/33 7s 207ms/step - accuracy: 0.8394 - loss: 0.3925 - val_accuracy: 0.7704 - val_loss: 0.5263 - learning_rate: 3.1250e-05
Epoch 25/50
33/33 7s 210ms/step - accuracy: 0.8245 - loss: 0.4087 - val_accuracy: 0.7704 - val_loss: 0.5312 - learning_rate: 3.1250e-05
Epoch 26/50
33/33 7s 209ms/step - accuracy: 0.8502 - loss: 0.3944 - val_accuracy: 0.7899 - val_loss: 0.4641 - learning_rate: 3.1250e-05
Epoch 27/50
33/33 7s 206ms/step - accuracy: 0.8667 - loss: 0.3630 - val_accuracy: 0.7899 - val_loss: 0.4650 - learning_rate: 3.1250e-05
Epoch 28/50
33/33 7s 212ms/step - accuracy: 0.8238 - loss: 0.3927 - val_accuracy: 0.7938 - val_loss: 0.4543 - learning_rate: 3.1250e-05
Epoch 29/50
33/33 7s 210ms/step - accuracy: 0.8281 - loss: 0.3801 - val_accuracy: 0.8016 - val_loss: 0.4476 - learning_rate: 3.1250e-05
Epoch 30/50
33/33 7s 207ms/step - accuracy: 0.8490 - loss: 0.3730 - val_accuracy: 0.8132 - val_loss: 0.4231 - learning_rate: 3.1250e-05
Epoch 31/50
33/33 7s 208ms/step - accuracy: 0.8522 - loss: 0.3835 - val_accuracy: 0.8249 - val_loss: 0.4292 - learning_rate: 3.1250e-05
Epoch 32/50
33/33 7s 209ms/step - accuracy: 0.8446 - loss: 0.3842 - val_accuracy: 0.8016 - val_loss: 0.4499 - learning_rate: 3.1250e-05
Epoch 33/50
33/33 7s 213ms/step - accuracy: 0.8359 - loss: 0.4148 - val_accuracy: 0.8132 - val_loss: 0.4167 - learning_rate: 3.1250e-05
Epoch 34/50
33/33 7s 210ms/step - accuracy: 0.8462 - loss: 0.3731 - val_accuracy: 0.8054 - val_loss: 0.4367 - learning_rate: 3.1250e-05
Epoch 35/50
33/33 7s 215ms/step - accuracy: 0.8420 - loss: 0.3817 - val_accuracy: 0.8171 - val_loss: 0.4263 - learning_rate: 3.1250e-05
Epoch 36/50
33/33 7s 207ms/step - accuracy: 0.8433 - loss: 0.3898 - val_accuracy: 0.8093 - val_loss: 0.4275 - learning_rate: 3.1250e-05
Epoch 37/50
33/33 8s 256ms/step - accuracy: 0.8439 - loss: 0.3793 - val_accuracy: 0.8210 - val_loss: 0.4259 - learning_rate: 1.5625e-05
Epoch 38/50
33/33 9s 263ms/step - accuracy: 0.8468 - loss: 0.3618 - val_accuracy: 0.8288 - val_loss: 0.4173 - learning_rate: 1.5625e-05
Epoch 39/50
33/33 9s 264ms/step - accuracy: 0.8143 - loss: 0.3974 - val_accuracy: 0.8327 - val_loss: 0.4141 - learning_rate: 1.5625e-05
Epoch 40/50
33/33 8s 233ms/step - accuracy: 0.8396 - loss: 0.3859 - val_accuracy: 0.8171 - val_loss: 0.4160 - learning_rate: 1.5625e-05

Epoch 41/50
33/33 7s 221ms/step - accuracy: 0.8620 - loss: 0.3517 - val_accuracy: 0.8288 - val_loss: 0.4143 - learning_rate: 1.5625e-05
Epoch 42/50
33/33 7s 218ms/step - accuracy: 0.8446 - loss: 0.3761 - val_accuracy: 0.8327 - val_loss: 0.4152 - learning_rate: 1.5625e-05
Epoch 43/50
33/33 7s 210ms/step - accuracy: 0.8603 - loss: 0.3681 - val_accuracy: 0.8249 - val_loss: 0.4138 - learning_rate: 1.0000e-05
Epoch 44/50
33/33 7s 212ms/step - accuracy: 0.8545 - loss: 0.3593 - val_accuracy: 0.8288 - val_loss: 0.4116 - learning_rate: 1.0000e-05
Epoch 45/50
33/33 7s 214ms/step - accuracy: 0.8366 - loss: 0.3784 - val_accuracy: 0.8288 - val_loss: 0.4097 - learning_rate: 1.0000e-05
Epoch 46/50
33/33 7s 212ms/step - accuracy: 0.8414 - loss: 0.3689 - val_accuracy: 0.8327 - val_loss: 0.4096 - learning_rate: 1.0000e-05
Epoch 47/50
33/33 7s 213ms/step - accuracy: 0.8626 - loss: 0.3538 - val_accuracy: 0.8327 - val_loss: 0.4102 - learning_rate: 1.0000e-05
Epoch 48/50
33/33 7s 214ms/step - accuracy: 0.8652 - loss: 0.3459 - val_accuracy: 0.8210 - val_loss: 0.4138 - learning_rate: 1.0000e-05
Epoch 49/50
33/33 8s 232ms/step - accuracy: 0.8523 - loss: 0.3621 - val_accuracy: 0.8327 - val_loss: 0.4120 - learning_rate: 1.0000e-05
Epoch 50/50
33/33 8s 245ms/step - accuracy: 0.8671 - loss: 0.3479 - val_accuracy: 0.8327 - val_loss: 0.4133 - learning_rate: 1.0000e-05



```
In [14]: # Plotting Loss
plt.plot(history.history['loss'])
```

```
plt.plot(history.history['val_loss'])
plt.title('Model Loss')
plt.ylabel('Loss')
plt.xlabel('Epoch')
plt.legend(['Train', 'Validation'], loc='upper left')
plt.show()
```



```
In [15]: # Prediction example (optional)
predictions = model.predict(X_test)
print(predictions[:5])
```

```
264/264 ━━━━━━━━ 11s 42ms/step
[[0.13961528]
 [0.29620218]
 [0.01766438]
 [0.9715492 ]
 [0.253062  ]]
```